


CPU microarchitecture regimes across the service pods (rag, tok320)



Pod classified by its largest top-down bucket (% of CPU pipeline slots). Below each: the hardware counter for that mechanism — MPKI = cache misses per 1000 instructions (L1-I\$ = instruction cache, LLC = last-level cache), or execution μ ops/cycle.

Retiring Frontend-bound Bad spec. Backend · memory Backend · execution