


vLLM is GPU-bound: the CPU just blocks per token

vLLM pod is CPU-idle ($\approx 1\%$)



...but switches scale with output (waiting on GPU)

