


Where request time goes (tok320)



- GPU (vLLM) 99.4%
- CPU (frontend) 0.58%