

Where each pod's CPU stalls — TMA top-down (rag, tok320)

