


Where request time goes (tok64)



- GPU (vLLM) 96.0%
- CPU (frontend) 3.95%