


Where request time goes (tok192)



- GPU (vLLM) 98.3%
- CPU (frontend) 1.66%